


1. Binary trees

- (a) (4 points) Let S be the set of binary trees (not necessarily search trees) over the set of keys $\{1, 2, 3, 4, 5, 6, 7\}$ that have in-order traversal $7, 4, 2, 5, 1, 3, 6$ and post-order traversal $7, 4, 5, 2, 6, 3, 1$. Answer the following questions.

1. What is the size of S ?

Answer: 1

2. Draw any one binary tree that belongs to S .



- (b) (5 points) Suppose you are given the pre-order and post-order traversals of a binary tree T over some set S of keys. Suppose $x, y \in S$. How will you determine whether x is an ancestor of y in T ? Briefly explain why your answer is correct.

x is ancestor of $y \iff x$ appears before y in pre-order and x appears after y in post-order.

\Rightarrow is obvious by def of pre and postorder

If y ancestor of x then y appears before x in preorder and y appears after x in post-order.

If y, x aren't ancestors/descendant of each other, let z be their lowest common ancestor.

preorder(z) and postorder(z) visit x, y in the same order.

2. (5 points) **String Processing**

Recall the pattern matching problem discussed in class and in the fourth assignment. Let Σ be a finite alphabet and x be a string over Σ . If p is a non-empty string over Σ , then we say that the pattern p matches x at index i if $x[i \dots (i + m - 1)] = p$, where m is the length of p . If p_1 and p_2 are non-empty strings over Σ , we say that the pattern $p_1?p_2$ matches x at index i if $x[i \dots (i + m_1 - 1)] = p_1$ and $x[(i + m_1 + 1) \dots (i + m_1 + m_2)] = p_2$, where m_1 and m_2 are the lengths of p_1 and p_2 respectively. In other words, '?' is a wildcard character which is not in Σ , and which matches every character in Σ .

Suppose you are given a text x and two patterns, p_1 and p_2 , which don't contain the wildcard. You have already computed L_1 and L_2 , which are sorted lists of indices at which p_1 and p_2 respectively match x . Write down an algorithm to compute L , the list of indices at which the pattern $p_1?p_2$ matches x . Your algorithm must run in time linear in the total size of L_1 and L_2 .

Idea: Subtract $m_1 + 1$ from each element of L_2 to get L_3 . Then find $L_1 \cap L_3$ using the "merge" step of merge-sort.

$L_3 \leftarrow \{ k - (m_1 + 1) \mid k \in L_2 \}$ (in sorted order)

$j_1 \leftarrow 0, j_2 \leftarrow 0$. $L \leftarrow$ empty list

while $j_1 < |L_1|$ and $j_2 < |L_3|$:

 if $L_1[j_1] = L_3[j_2]$

 Append $L_1[j_1]$ to L .

$j_1 \leftarrow j_1 + 1, j_2 \leftarrow j_2 + 1$

 else if $L_1[j_1] < L_3[j_2]$

$j_1 \leftarrow j_1 + 1$

 else

$j_2 \leftarrow j_2 + 1$

return L .

3. (4 points) **(2,4)-trees**

Recall the procedure for deleting a key from a (2,4)-tree: if we find the required key in a non-leaf node, then we exchange it with its successor and then delete it. But how do we find its successor in the first place? Your job is to write a **python** code to find the successor of a key in a (2,4)-tree that resides in a non-leaf node. A node in our (2,4)-tree has the following attributes.

- **key**: a list of 3 objects. If the node contains k keys (where $1 \leq k \leq 3$), then these keys are `key[0], ..., key[k-1]`, while `key[k], ..., key[2]` are all `None`.
- **child**: a list of 4 node-references. If the node has d children (where $0 \leq d \leq 4$), then the references of these children are `child[0], ..., child[d-1]`, while `child[d], ..., child[4]` all `None`.
- **parent**: this is `None` if the node is the root of the tree, else the reference to the parent node.

Complete the following function which takes a reference `p` of a non-leaf node and an integer `i` as input parameters, and returns a pair `(q, j)`, where `q` is a node-reference and `j` is an integer such that `q.key[j]` is the successor of `p.key[i]` in the (2,4)-tree. The function must run in time $O(\log n)$, where n is the number of keys in the (2,4)-tree.

```
def findSuccessor(p, i):
```

```
    q = p.child[i+1]
    while q.child[0] is not None:
        q = q.child[0]
    j = 0
    return (q, j)
```


4. Undirected Graphs

Recall that we defined the distance between two vertices of an undirected graph to be the length of a shortest path between those vertices (for example, the distance between a vertex and itself is 0, the distance between adjacent vertices is 1, and so on). The *diameter* of a graph is defined as the maximum, over all pairs u, v of its vertices, of the distance between u and v .

Let G be a connected undirected graph. Perform a breadth-first-traversal of G starting from some vertex s , and let T be the resulting breadth-first-traversal-tree. Suppose T has exactly $h + 1$ levels $L_0 \dots, L_h$, where $L_0 = \{s\}$.

- (a) (2 points) Prove that the diameter of G is at least h .

A vertex u in L_h is at distance h from s
 $\therefore \max_{u, v \in V} \text{dist}(u, v) \geq \text{dist}(s, u) = h.$

- (b) (4 points) Prove that the diameter of G is at most $2h$.

For any two vertices u, v ,
 $\text{dist}(u, v) \leq \text{dist}(s, u) + \text{dist}(s, v)$
 \uparrow
 Δ inequality $\leq h + h = 2h$
 \uparrow
 A vertex x belongs to L_i iff
 $\text{dist}(s, x) = i.$

5. Directed Acyclic Graphs

Recall the out-adjacency list representation of directed graphs, where we have one linked list for every vertex v , where the linked list contains the out-neighbors of v . Given a directed acyclic graph G in its out-adjacency list representation and a vertex t of G , we would like to compute, for every vertex v of G , the number of directed paths from v to t . Note that the number such paths could be exponential in the number of vertices, so a brute-force counting is too inefficient.

- (a) (3 points) Complete the following algorithm to compute an array `pathCount` indexed by the vertex set of G so that at the end of the run of the algorithm, `pathCount[v]` equals the number of directed paths from v to t in G . Each blank must be filled up with an expression that can be computed in $O(1)$ time.

`findPathCounts(G, t)`

1. Compute a topological sort of G .

2. For each vertex v in the reverse of the topological sort order:

2.1. If $v=t$ then `pathCount[v] ← 1`, else `pathCount[v] ← 0`.

2.2. For each out-neighbor u of v : `pathCount[v] ← pathCount[v] + pathCount[u]`

3. Return `pathCount`.

- (b) (4 points) Prove that the above algorithm runs in time $O(n + m)$, where n is the number of vertices and m is the number of edges in G .

Topological sort $\rightarrow O(m+n) \rightarrow$ proved in class
 Step 2.1 : once for every vertex $\therefore O(n)$
 Step 2.2 : $O(d_v)$ for vertex v , where d_v is out-degree of v .
 \therefore Total time for 2.2. $\therefore O(\sum_v d_v) = O(m)$

Name: _____

Entry number: _____

COL106

Major Exam

Duration: 2 hours

- (c) (4 points) Why is it necessary to iterate over the vertex set in reverse topological sort order in the algorithm of part (a)? Explain briefly.

We are using the fact: that

$$\# \text{ } v \text{ to } t \text{ paths} = \sum_{\substack{u: \text{ outnbr} \\ \text{ of } v}} \# \text{ } u \text{ to } t \text{ paths} \quad \text{if } v \neq t.$$

Reverse topological sort ensures that when we compute the LHS above, each term on the RHS is computed already.